



Training dynamics from 3730 complete epoch logs (NEU-CLS)

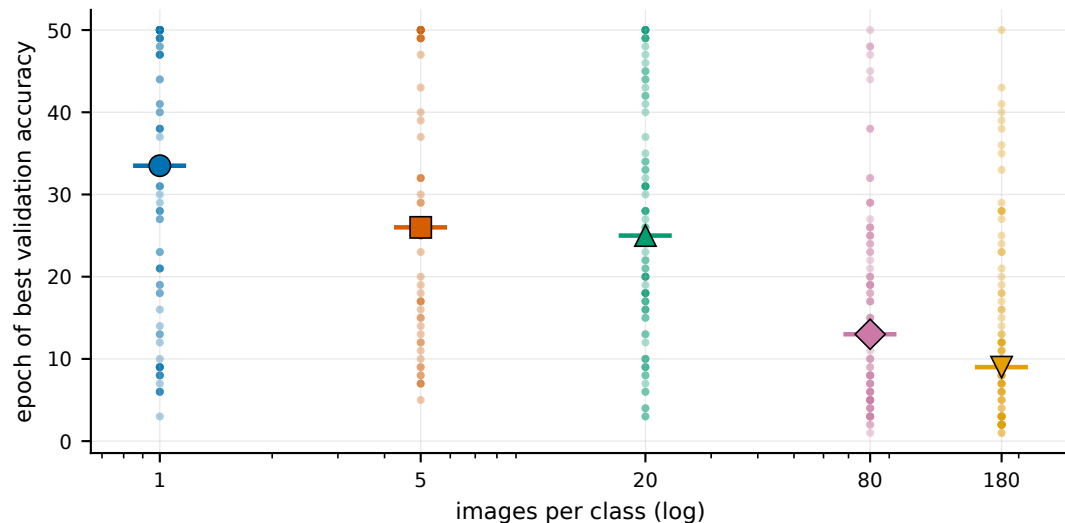
(a) Validation curves converge within ~15 epochs at every budget



(b) Generalisation gap closes as data grows (memorisation is a budget effect)



(c) Best epoch arrives earlier as data grows (median 34 -> 9 epochs as n goes 1 -> 180)



(d) The pre-dedup sweep is still usable
n=1265 matched runs, mean |diff| = 1.41 pp

